Amrita School of Computing
Amrita Vishwa Vidyapeetham
Amritapuri, India

November 14, 2026

[Programme Chairs]
[conference name]

Dear [Programme Chairs],

We submit "The Agent Holds the Key: Auditing What LLM Agents Learn Through Redaction Tools" to *[conference name]*.

LLM coding agents read workspace text through tools, and agents are increasingly given local redaction tools through the Model Context Protocol. The paper asks what an agent learns through such a tool. In a pre-registered study of an open-source on-device guardrail, a differential egress audit found that a reply documented as content-free returned person names through detector metadata; we fixed it and verified the fix with the same audit and a property test. We also measure a regex-gated neural name tier on real user-generated text, show that batched inference re-opens a chunking leak, and report the per-call cost of the guardrail, including a transport stall that the measurement exposed.

The manuscript is not under consideration elsewhere, and all authors approve the submission. A separate manuscript by the same authors, "Chunking Without Leaking: Provably Equivalent, Bounded-Memory Redaction of Personal Data and Credentials", is under review at a journal. It analyses the streaming engine and the pattern tiers, which this paper treats as a fixed baseline. The two papers ask different questions on different corpora and share no result; we will gladly send it to you. An AI assistant was used to implement the evaluation, run the experiments and draft the text, as the manuscript states; the authors take full responsibility for the content. The authors declare no competing interests.

Yours sincerely,

Parvathy R (corresponding author), on
behalf of all authors
Amrita School of Computing, Amrita
Vishwa Vidyapeetham