[MISSING: Department, Institution]
[MISSING: City, Country]

October 11, 2026

The Editors
Computers & Security
Elsevier

Dear Editors,

Please consider the enclosed manuscript, "Bounded-memory streaming redaction of personal data and credentials with a whole-document equivalence guarantee," for publication in *Computers & Security*.

Tools that remove personal data and credentials from logs and exports are widely deployed, but the artefacts they process are routinely larger than memory, so they must be processed in chunks. We could find no system in the redaction, de-identification or secret-scanning literature that states any relationship between its chunked output and its whole-document output. The gap matters operationally: a value straddling a chunk boundary can be written out in the clear, and nothing in the output signals that it happened.

The paper closes that gap. Given bounds on span length and context that each detector declares and that regression tests enforce, we prove the streamed output equals the whole-document output, the parallel output equals the sequential output byte for byte, and working memory is independent of input size. We then measure what this costs and what it buys, on held-out test splits against Microsoft Presidio in three configurations and against gitleaks and TruffleHog, with bootstrap confidence intervals, a component ablation, and a memory scaling curve from 8 MiB to 2 GiB.

We have tried to report the evaluation as a referee would want to see it. Where we lose, we say so: free-text recall without a named-entity model is low, and against Presidio's larger model our entity coverage is not significantly better. One corpus was inspected during rule development, and rather than drop it we report pre-exposure and post-exposure results side by side and bound the contamination. All code, split manifests, the evaluation harness and the raw results underlying every number are public.

[MISSING: Keep the next paragraph unless the FTNCT'09 paper was never submitted anywhere; if it was submitted, accepted or published, complete it and attach a copy.] An earlier version of part of this work appeared as [MISSING: full citation] at [MISSING: venue, date]. A copy is enclosed. That paper presented the streaming architecture, the equivalence theorems, parallel seam safety and the calibrated ensemble. The present article adds: evaluation on held-out test splits, which the conference paper did not report; secret-scanner baselines (gitleaks, TruffleHog); a component ablation; a contamination analysis of RedactionBench; a peak-RSS memory curve across a 256-fold range of input sizes; benign-shape suppression in the entropy tier, which reduced the false-alarm rate on generated non-secrets from 0.794 to 0.010; and parity results for the native core. We estimate the new technical content at [MISSING: X]% of the article.

The manuscript is not under consideration elsewhere, and all authors have approved the submission. We disclose in the Methods section that an AI assistant was used to implement evaluation scripts and to draft the manuscript; the authors reviewed every result and every sentence and

are accountable for the content.

Yours sincerely,

Mounish Senisetty
[MISSING: affiliation]
senisettymounish@gmail.com